

Security: How systems behave in the presence of an adversary

- "Defense in depth": Swiss cheese model, where multiple layers of defense cover potential "holes"
- Computer Fraud and Abuse Act (CFAA)

CRYPTOGRAPHY: Communicate securely in presence of an adversary — adversary typically modeled as MITM — can see, modify, and forge messages between A, B

3 GOALS (CIA):

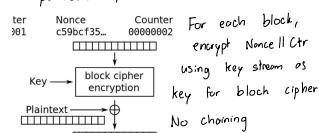
- Confidentiality: Mallory can't see what A, B are saying to each other.
- Integrity: Mallory can't modify msgs without being detected.
- Authenticity: Alice can confirm that a message really came from Bob, and vice-versa. May be a modified msg, but really from Bob.

ONE-TIME PADS / STREAM CIPHERS

- achieve CONFIDENTIALITY
- OTP: Alice and Bob share secret, long string of bits K
 - $\rightarrow C_i = p_i \oplus k_i \Rightarrow p_i = c_i \oplus k_i$
 - \rightarrow information - theoretically secure
 - \rightarrow CANNOT REUSE PAD — leak information
($a \oplus k$) \oplus ($b \oplus k$) = $a \oplus b$
 - \rightarrow usually impractical - how to exchange k ?
- Stream ciphers: Use PRG for OTP
 - \rightarrow Alice and Bob choose PRG family G , w/ secret key K (g, k)
 - \rightarrow Need new key every time, never reuse PRG output bits — but provably secure!
 - \rightarrow Can't prove secure PRGs exist \therefore
 - \rightarrow Broken: RC4 ; Good: ChaCha20

COUNTER (CTR) Cipher Mode

- Turn a block cipher into stream
- Using a PRG, generate stream using K
- Have a counter (incremented per block)
- Have unique nonce per message
- Benefits: don't need padding, efficient parallelism / random access



NEVER REUSE NONCE FOR SAME K

DIFFE-HELLMAN Key Exchange

- Publicly known params:
 - \rightarrow Large prime p (2048 bits)
 - \rightarrow Generator g for p (usually small)
- Alice chooses secret key a , computes and sends $g^a \bmod p = A$
- Bob chooses b , sends $g^b \bmod p = B$
- $S = B^a \bmod p = A^b \bmod p = g^{ab} \bmod p$
- S is secret key!
- VULNERABILITY: MITM attack succeeds
 - \rightarrow Alice and Bob each compute secret key $s = k$ / Mallory — she has to intercept + transform every msg, but can read

- Defenses:
 - \rightarrow Trust on first use, detect later (SSH)
 - \rightarrow Communicate Separately to verify same key
 - \rightarrow rely on proximity to hope no MITM
 - \rightarrow digital signature! HTTPS does this

FORWARD SECRECY

- if key stolen, don't want to crack previous messages
- Use D-H to generate sess keys
- Authenticate session keys w/ separate long-term key
- If long-term key leaked, past still Secure

MESSAGE INTEGRITY

- Idea: Have some verifier V for each message m . Verifier should be computable by Alice + Bob, but not Mallory (at least not easily)



Bob accepts iff $f(m') = V$

\Rightarrow ideally, this can only happen when $m = m'$, $V = V'$ as f unknown to Mallory

Want f to "look random"

- ideally, a random function (too infeasible)
- so, we use PRF like HMAC-SHA256
- can't prove this is a PRF (would imply $P = NP$) but assumed
- For m , compute $V = \text{HMAC-SHA256}(m)$, with exact function in fn family determined by secret key K (known to Alice + Bob, but not Mallory)

HASH FUNCTIONS

- **PRF**: Mallory should not be able to distinguish PRF from known function family of size 2^n from truly random function in polynomial time
 - \rightarrow "random": deterministic but random
 - $\rightarrow n = \#$ of bits in key
 - \rightarrow could always do in exp time by checking against each possible $K \in \{2^n\}$
- Hash functions act as PRFs
 - no keys, just a fixed function
 - idea: prepnd key to msg and hash to get V — this fails due to length-extension attack
 - MD-5, SHA-1 broken
 - SHA-256, SHA-512, SHA-3 secure

HMAC

MAC is method authentication code

$$V = \text{HMAC}_K(m) = H(K \oplus C_1 \| H(K \oplus C_2 \| m))$$

H : any hash function \oplus : XOR

C_1, C_2 : publicly known constants

BLOCK CIPHERS

- Encrypt each n -bit block w/ some (reusable) key K

$$E_K(p) \rightarrow c; D_K(c) = E_K^{-1}(c) \cdot p$$

- \rightarrow K selects what function to use; there are $(2^n)^n$ functions (2^n strings, $(2^n)^n$ ordering, each ordering corresponds to a $2^n \leftrightarrow 2^n$ bijection)
- \rightarrow This is a Pseudorandom Permutation b/c it is bijective; a PRF compresses.

CIPHER Modes

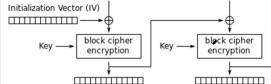
- Given black-box block cipher, how do we chain it to encrypt msg?
- ECB is flawed!:
 - encrypting each block independently with same key leaks info

Cipher Block Chaining (CBC):

- "Chain" ciphertexts to obscure
- Need initialization vector IV (random)

- sent unencrypted
- must be diff each msg

Initialization Vector (IV):



- Problems: can't encrypt in parallel; have to send IV w/ ciphertext each time

RSA (Rivest, Shamir, Adleman)

- Public-key cryptography — previous paradigm are 1-1 comm. What if you want to publish data + show integrity?

- Choose large primes P, Q : $N = PQ$

$$\rightarrow$$
 Pick small e , s.t. relatively prime to $\phi(N) = (P-1)(Q-1)$

$$\rightarrow$$
 Compute d s.t. $ed \equiv 1 \pmod{\phi(N)}$

• Public key e , modulus N

• Anyone can encrypt: $C = m^e \bmod N$

• To decrypt: $m = C^d \bmod N$ ($m < N$)

• RSA Signatures:

Sign m w/ $s = m^d \bmod N$
Verify: $s^e \bmod N \stackrel{?}{=} m$

Must be for FIXED msg

• Drawbacks: AES 1000x faster; 3 subtle mathematical attacks

- Bleichenbacher's Attack: If e is small (say 3), $m \in N^{\frac{1}{e}}$, can just take cube root of c
- Key gen failures
 - \rightarrow Mallory can trick Bob into signing sensitive msg
- Forge arbitrary msg signatures

Forging signatures on specific messages:

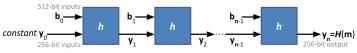
Suppose Bob will sign any message Mallory provides that he deems non-sensitive: Mallory can trick Bob into signing sensitive msg: Mallory computes: $m := r^e \bmod N$ for random r . Bob signs m : $s := m^d \bmod N = (r^e)^d \bmod N$. Then, Mallory finds signature $s' = (m')^d \bmod N$ by computing $s' = (r^e)^d \bmod N = (r^d)^e \bmod N$

LENGTH - EXTENSION ATTACKS

- Many hash functions, like SHA-256, use Merkle-Damgard construction to hash arbitrary-length inputs to 256 bits.
 - \rightarrow "random": deterministic but random
 - $\rightarrow n = \#$ of bits in key
 - \rightarrow could always do in exp time by checking against each possible $K \in \{2^n\}$
- If $V = f(m) = \text{SHA-256}(K \| m)$, this is vulnerable!
- Given $V = f(m) = \text{SHA-256}(K \| m)$, for unknown m , attacker can calculate constant b_0 256-bit inputs b_1, b_2, \dots, b_{n-1} b_n $b_0 \oplus b_1 \oplus b_2 \oplus \dots \oplus b_{n-1} \oplus b_n = V$, where $m' = m \| 11$ padding || arbitrary
- use digest V as y_0 for rest of message — as key already prepended to m , everything else (data b_i , padding scheme) is publicly known.
- Final m' that hashes to V' will be:
 - \rightarrow key || original msg || orig padding || new msg || new padding

M-D construction:

- Have some compression func h
- Input: b_0 (512 bit), y_0 (256 bit)
- Output: 256 bits
- Split m into $b_0, b_1, \dots, b_{n-1}, b_n$ with padding as required — ALWAYS add some padding
- Have some publicly known y_0 (random 256-bit)



GENERATING RANDOMNESS

- Need to select a truly random key
- true randomness is hard; we use PRG
- PRG: Function g that isn't distinguishable from truly random function in poly-time
 - \rightarrow generates a STREAM of bits
- Given PRF f w/ key K , $g(0) = f_K(0) \parallel f_K(1) \parallel f_K(2) \dots$ is PRG (random stream of bits)
- RNG is a good attack target — knowing PRG is VERY POWERFUL
- Eg: NSA backdoored Dual-ECC DRBG (2006)

AUTHENTICATED ENCRYPTION

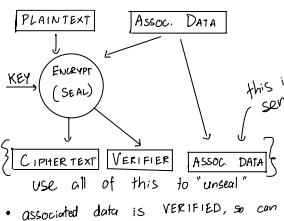
- APPROACH 1: Compose encryption + MAC
- APPROACH 2: Build primitive to do both
- Basically, encrypt fn can return "fail" now

HOW TO COMPOSE MAC + ENCRYPT

- CRYPTOGRAPHIC Doom: If you perform any cryptographic operations on a message without verifying MAC FIRST, you're screwed
- Encrypt-and-MAC: $E(p) \parallel MAC(p)$
- MAC-then-encrypt: $E(p \parallel MAC(p))$
 - \rightarrow Vulnerable to padding oracle attack
- Encrypt-then-MAC: $E(p) \parallel MAC(E(p))$
 - \rightarrow THIS IS SAFEST

AUTHENTICATED ENCRYPTION w/ Assoc Data

- AEAD — use unencrypted assoc data



- associated data is VERIFIED, so can check MAC first

- AES-GCM (Galois Counter Mode)
- ChaCha-Poly1305, common on mobile
- AES-GCM: Seal function consists of GTR encryption for plaintext, and Galois Field MAC of ciphertext + assoc data
- Same key used for both (non-generic composition)
- NEVER REUSE NONCE

HTTP Protocol

- Sequence of requests and responses for retrieving data
- Client sends: Method (GET / POST), path query, headers
- DOM: Document Object Model API to allow JS to modify HTML
- Pages can include scripts / load them from external resources - escape certain chars - < > & " '

URLs

Scheme: //host: port /path ? query # fragment

- scheme: http, https, etc.
- host: IP addr / domain name
- port: TCP port (default 443 for https, 80 http)
- path: which resource on server?
- query: parameter string passed to server
 - often key-value pairs, could be arbitrary
- fragment: not visible to client (what section of page?)

COOKIES

- Cookies use a different scope than DOM — this leads to problems
- Scoped on ([scheme], domain*)
- Secure attribute: prevent cookies being sent over http
- HttpOnly attr: prevent cookies from being read by JS running in origin

- Site can SET cookie for own domain or any parent domain as long as parent domain is not public suffix (eg .com, .gov, .edu)
- You don't know who set a cookie when you receive it
- Site can READ cookie for own domain, parent domains. Scripts can read cookie for any path on domain — but diff path cookies not sent by default

SAME ORIGIN Policy (SOP)

- Separates content into origins; restricts data flows between them. Eg: attacker.com can't load Gmail in an iFrame and read your messages
- Origin: scheme://domain:port
- What's isolated?
 - Cookies, DOM storage, JS namespace, local hardware perms
- ALLOWED: cross-origin writes (links, redirects, form submissions), cross-origin embedding (images, videos, iFrames, etc.)
- CORS: server can allow cross-origin requests from specified origins

Estimating what's feasible to compute?

$2^{64} \approx 10^{19}$

$2^{128} \approx 10^{38}$

1 year =

3×10^7 s

GPU mining: $\approx 10^8$ SHA-256/s

GPU mining: $\approx 10^{11}$ SHA-256/s

ASIC mining: $\approx 10^{14}$ SHA-256/s

Bitcoin miners globally: 10^{20} SHA-256 blocks/s

CSRF ATTACKS

- Core idea: Cookies being sent are determined by domain of resource being requested
- if a page on attacker.com requests image from bank.com, bank cookies sent (unless sameSite attr. set, etc.)
- Malicious site can send requests to other sites - auth cookies will be sent w/ it
- eg. send POST req. to bank.com to transfer money; login to site under attacker act to monitor activity
- Cookie-based authentication ALONE BAD

CSRF DEFENSES

- Referrer Validation: HTTP request contains URL of page making request (however, not always sent! Doesn't always work)
- Secret Token Validation: Pages on actual site include secret token (Csrf token) in all requests - helps block cross-site requests in general (Session-dependent token). Attacker cannot retrieve due to SOP
- Same Site Cookies: SameSite attribute prevents cookie from being sent cross-site → Strut: not sent in ANY context, even when following regular link → Lax: Sent on cross-site links, but not subrequests (usually want this)

XSS DEFENSES

- Validation + Escaping: Check all user input robustly to prevent any scripts → also encode all special characters in output → adopt positive security policy - specify what IS allowed, rather than what isn't (less likely to have loopholes)
 - Content Security Policy:
- http header, tightly specifies what scripts allowed to execute (what domains scripts must come from) - inline scripts banned. More modern approach

PROJECT 2: FOR SQLI, TERMINATE STRING W/ ' AND THEN DO OR 1=1 # TO COMMENT OUT

For XSS, can insert between <script> tag or onerror of img tags, etc. Insert into URL s.t. it will be put into HTML somehow (echo user inputs). Wait for document to load. This is vulnerable!

For CSRF, can have hidden code on attacker's website that logs in/ submits other req. to other websites. Use POST reqs w/ FORMS, reduced to frames. If there are XSS vulnerabilities, can bypass CSRF tokens by setting/getting cookies.

PROJECT 1: Vulnerability exists when using pure hash functions as MACs. Can do length extension to hash some longer message, even if key unknown.

bytes frommer(...)|SHA256(m).hmac

```
$ openssl dgst -sha256 file1
$ fsockcall -o file2 file2 -p prefix
→ generates MD5 collisions w/ some prefix and diff ending
→ Hashes will be same if some suffix added (Merkt-Damgord)
```

BLEICHENBACHER:
padding not checked,
so less FPs -
this makes m³ < n,
so ciphertext is
easy to decode;
Can forge RSA
signature

PHYSICAL LAYER (L1): "Encrypt @ higher"

- Easy to tap, jam, intercept. Historically poor encryption - WPA, 20, 30, 40 all broken
- LINK LAYER (L2): How do I get to next hop?
- Assume: some "physical" link b/w nodes
- Ethernet/WiFi: Send packets (frames) to hosts on local network (LAN)
- Devices ship w/ globally-unique 48-bit MAC address (but can change arbitrarily)
- EtherType: L3 protocol spec
- Ethernet Switches: each port connects to host / other switch. LEARNS MACs on each port eventually to route - broadcast to all if unknown
- These do not scale! We need IP for L3
- Plaintext; packets could be corrupted, lost, etc.

SQL INJECTIONS

- Website does not sanitize inputs to SQL statements - can have data interpreted as code (DANGEROUS)
- Typical SQL statement:
SELECT * FROM users WHERE username = '\$user' AND password = '\$pwd'
- What if \$pwd ends SQL stmt and does other stuff? Ex:
pwd = ' OR 1=1 -- (selects all rows from DB)
- DEFENSE: Don't build SQL cmds yourself - use parametrized SQL statements, where query and arguments are passed separately
- XSS (Cross-site Scripting)
- Type of injection - data interpreted as code by insecure site
- Client-side attack

REFLECTED XSS

- Vulnerability when site echoes inputs back to user w/o cleaning
- Could put code in query of malicious URL - site could accidentally put it in HTML (thinking it's text), and it gets executed
- DANGER: Even one page having vulnerability exposes whole site, as script not blocked by SOP any longer

STORED XSS

- Vulnerability: Site stores + displays unsanitized user content to other people
- Attacker can upload a script as part of some content. This content could be accessed by other users, at which point script is executed on their browsers!
- Samy Worm: 2005. Uploaded Script to MySpace that replicated in user profiles when they viewed his profile. Propagated across site.

NETWORK LAYER (L3): Packets to final dest

- IP (v4, v6). Routing + Fragmentation - no security; pxtl packets could be lost

IP header on every packet: source addr (unverified), dest addr, other stuff

v4 addr: 32 bits, exhausted

v6 addr: 128 bits, AA:BB::XX:YY

→ :: means all 0s

IP has network prefix, host suffix

→ CIDR notation: how many bits in prefix?

→ 192.168. ... (24) OR a mask

Same prefix = same LAN - just go over ethernet. Else, send to router (network gateway)

PRIVATE IPv4s:

10.0.0.0 → 10.255.255.255

192.168.0.0 → 192.168.0.255

Residential ISPs only get 1 addr

- Home routers assign each device address from private range; rewrite packets in both directions

localhost: 127.0.0.1

THREAT MODEL:

- Off-path: talk to hosts, but can't see packets (most attackers @ first)
- On-path: See + add packets, but can't modify or block - ethernet passive tap, for instance
- In-path: MITM, full power

Packet Capture: Hosts can log all packets arriving on network interface (pcap format, tools like Wireshark).

mtr, traceroute to see how your packets travel

TLS (TRANSPORT LAYER SECURITY)

- Protocol stacked on top of TCP, which is a plaintext communication protocol. → TCP provides "phone-call" semantics
- HTTP → TLS → TCP → Internet
- Threat model: malicious network
 - Govt. censorship, data harvesting, compromised routers, route hijacking, etc.
- PROVIDES: Confidentiality + Integrity for data while in transit. Client can also authenticate server's identity
- Does NOT Provide: Malware resistance, tracking, metadata analysis (ie which sites you visited), DDoS, etc.
- TLS 1.3 (2018): Major redesign
- Confidentiality + Integrity thru AEAD ciphers
- Authentication w/ certificates + public key crypto

TLS CERTIFICATES

- Server signs hash of handshake w/ private key and provides public key. How do we know public key is legit??
- Server provides a certificate w/ identity + public key. This is signed by some Certificate Authority (CA)
 - public keys of CAs hardcoded into browsers (root CAs)
 - TLS uses X.509 certificates
- Often a chain of certificates leading back to some root CA
 - root CA has self-signed certificate w/ public key - this means that even if root key stolen, existing certs still legit b/c signed with public key who's private key not known
- Obtaining a certificate:
 - Server proves identity thru email, http, DNS
 - ACME: open protocol that automates this

PUBLIC KEY INFRASTRUCTURE (PKI)

- Maintained by CAs (100s exist), browser developers, CA/Browser forum
- Revocation lists exist if CA revokes
- Browser shows warning if cert expired
- HTTPS now ubiquitous - negative indicators now for http

ADDRESS RESOLUTION PROTOCOL (ARP)

- map IP addresses to MAC addresses (LAN)
- Host that needs MAC for IP broadcasts ARP packet to LAN - whoever has it replies
- Host caches this mapping

ARP SPOOFING

Anyone on LAN can claim to have that MAC to both user and router

- become MITM. Typically no prevention for this - must analyze @ network-level

ROUTER

Device w/ MULTIPLE link layers, each on diff network. Forwards packets to neighbor to get it closer to destination.

- Border Gateway Protocol (BGP): ISPs use to announce network prefixes, connectivity. Routers compute route tables from this.
- no authentication! ISP can announce someone else's network. BGP Hijacking is common!
- RPKI: sign records - slow adoption

TRANSPORT LAYER (L4): Add ports, encryption

- Application on server on top of bare packets identified by 16-bit port # - allow concurrent connections. 0-1023 usually require superuser priv.
- 32768 - 65535 EPHEMERAL - use for source ports
- HTTP: 80 + HTTPS: 443 + SMTP: 25 + SSH: 22

USER DATAGRAM PROTOCOL (UDP): Lightweight wrapper on IP (L4)

- Unreliable data transfer, but very fast!
- Used by real-time communications apps, DNS, QUIC
- Vulnerable to impersonation by off-path attacker

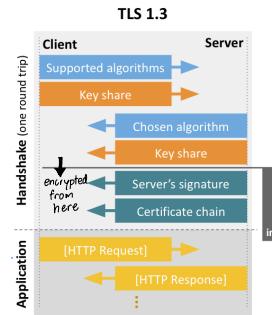
TRANSMISSION CONTROL PROTOCOL (TCP):

- Connection-oriented, reliable byte stream, congestion control. NO STRONG SECURITY

Sequence #s, acknowledgement #s to reassemble

TLS HANDSHAKE

- Negotiate Crypto algorithms
 - find mutually-supported ones
 - Key Exchange Algo (eg. DHE)
 - Signature Algo (eg. RSA w/ SHA-256)
 - Symmetric Crypto Algo (AEAD like AES-GCM)
- Establish Shared Secret
 - Diffie-Hellman for forward secrecy
- Authenticate Server
 - Server Signs hash of all previous messages w/ private key
 - sends certificate chain to authenticate public key - client checks
 - if public key not verified, anyone can sign message and give "correct" public key.



Client guesses which key-exchange algo server will pick - easy if comm before. Reduces expected # of round-trips.

FOOLING USERS

- Homographs: Attackers can buy look-alike domain names
- IDN homographs: different unicode chars look the same; can get perfect looking match for domain name in browser → only heuristics for defense
- Stripping: Some browsers default to http (SSL strip)
- MITM connects w/ server using https, but sends to user over http - now can read user data
- Mitigation: HSTS. Server sends header on first visit to only use https for x seconds. HSTS preload lists shipped w/ browser for first visit protection → fail if https not available
- Phishing: https cannot prevent - browsers use ML to detect

BUGS IN TLS

- TLS complicated, hard to code
 - Eg: Apple accidentally skipped cert checking for a year
 - Null prefix attack: diff in Pascal-C-style strings
 - Memory leaks leak random data
 - Mitigate using formal verification - miTLS is formally-verified lib

SERVER VULNERABILITIES

- TLS servers must protect private keys - always have isolation, access control. Contact CA for cert revocation if keys leaked

TCP HANDSHAKE: TCP carries data in segments

- Session starts w/ 3-way handshake - Client and server pick init sequence number, acknowledge each others
- On-path TCP attacks are trivial
- Off-path attacker can't do handshake without guessing server's seq. number - > 2^32, so some security props

APPLICATION LAYER (L5):

- Eg: SSH, DNS, SMTP, HTTP, etc.
- DNS: Delegatable, hierarchical database used to resolve hostnames to IP addresses
- control over regions of namespace (unidirectional delegation to AUTHORITY servers (DNS zone))
- 13 root authoritative servers exist; these delegate to servers for .com, .edu, etc.
- Cache responses for TTL (time to live)
- DNS Hijacking: Attacker hacks router so it points requests to attacker's DNS server; now, attacker can give whatever IP address for any domain name. DNS cache poison.
- Hosts File Hijacking: OS has DNS override file - can try and change this.
- DNS monitoring/blocking: Plaintiff DNS allows network to see what websites you are visiting (and potentially stop you)
- DEFENSES: DNSSEC adds authentication to DNS responses using TLS-like PKI. Issues: slow adoption.

OFF-PATH DNS CACHE POISONING (2008):

- Victim visits attacker's site, which executes some JS that sends bulk DNS requests to [??].bank.com
- Attacker can spoof responses of queries by guessing IP block query is at, being close to victim. Can temporarily redirect domain to whatever IP addr.
- Can increase randomization by randomizing source IP addr, randomizing ciphertext in query

DOS ATTACKS: Overwhelm host/network w/ traffic

- Assymetry: when small requests have large response
- TCP-SYN flooding: Space allocated before handshake complete. Fixed by making server send more MAC of client's, so no space req.
- Amplification: attacker queries server's w/ victim's IP as the same, inundating victim w/ responses
- Botnets: Many devices controlled by 1 person
- ISPs should drop packets w/ spoofed IPs to prevent
- CDNs like Cloudflare can help
- DNS over-TLS: DNS over-HTTPS operates DNS over encrypted channel (TLS instead of plaintext UDP)

